

# Gayathri\_Gopalan\_DS\_Challenge\_Shopify

September 21, 2021

## 0.1 Winter 2022 Data Science Intern Challenge

- 
- 

Question 1: Given some sample data, write a program to answer the following: [click here to access the required data set](#)

```
[14]: import pandas as pd
```

```
[15]: df= pd.read_csv('2019 Winter Data Science Intern Challenge Data Set - Sheet1.
↳csv')
```

```
[16]: df.head(4)
```

```
[16]:
```

	order_id	shop_id	user_id	order_amount	total_items	payment_method	\
0	1	53	746	224	2	cash	
1	2	92	925	90	1	cash	
2	3	44	861	144	1	cash	
3	4	18	935	156	1	credit_card	

	created_at
0	2017-03-13 12:36:56
1	2017-03-03 17:38:52
2	2017-03-14 4:23:56
3	2017-03-26 12:43:37

On Shopify, we have exactly 100 sneaker shops, and each of these shops sells only one model of shoe. We want to do some analysis of the average order value (AOV). When we look at orders data over a 30 day window, we naively calculate an AOV of \$3145.13. Given that we know these shops are selling sneakers, a relatively affordable item, something seems wrong with our analysis.

**0.1.1 a. Think about what could be going wrong with our calculation. Think about a better way to evaluate this data.**

AOV (Average Order Value)- It is the average amount spent each time a customer places an order of the shoe Since Sneakers are relatively affordable, and also that an individual cannot spend an

average of \$3145 on an order.

Current AOV calculation =  $(x_1+x_2+x_3+\dots+x_n)/n$

This calculation of the average of order amount is wrong. Since it has some of the value that have huge numbers. Let's take the median and standard deviation of this column.

Metric	Value
Median	284
Std Dev	41,282.54
Average	3145.128


```
[30]: df.describe()
```

```
[30]:
```

	order_id	shop_id	user_id	order_amount	total_items
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000
mean	2500.500000	50.078800	849.092400	3145.128000	8.78720
std	1443.520003	29.006118	87.798982	41282.539349	116.32032
min	1.000000	1.000000	607.000000	90.000000	1.00000
25%	1250.750000	24.000000	775.000000	163.000000	1.00000
50%	2500.500000	50.000000	849.000000	284.000000	2.00000
75%	3750.250000	75.000000	925.000000	390.000000	3.00000
max	5000.000000	100.000000	999.000000	704000.000000	2000.00000

We can see that the standard deviation is too high. It implies that the dispersion of data is too much which could be due to outliers.

```
[36]: ax2 = df.plot.  
      ↪scatter(x='order_id',y='order_amount',colormap='viridis',title="OrderID Vs_  
      ↪OrderAmount")
```



This proves that the outliers (in range of 200k-700k) is causing the average to an unrealistic number. Now let's try to remove the records which have unexpected order amounts (>100000).

```
[19]: df2=df[df['order_amount'] < 100000]
```

```
[20]: df2.describe()
```

```
[20]:
```

	order_id	shop_id	user_id	order_amount	total_items
count	4981.000000	4981.000000	4981.000000	4981.000000	4981.000000
mean	2501.425216	50.095162	849.915880	702.748444	1.992773
std	1443.152284	29.052172	86.816531	4627.726634	0.981326
min	1.000000	1.000000	700.000000	90.000000	1.000000
25%	1251.000000	24.000000	776.000000	163.000000	1.000000
50%	2503.000000	50.000000	850.000000	284.000000	2.000000
75%	3751.000000	75.000000	925.000000	390.000000	3.000000
max	5000.000000	100.000000	999.000000	77175.000000	8.000000

The std dev has reduced, but is still higher. Lets also take a look at max value for the 4th quartile order amt is 77175.00. Let's further remove the orders greater than 10k

```
[21]: df2[df2['order_amount'] < 10000].describe()
```

```
[21]:
```

	order_id	shop_id	user_id	order_amount	total_items
count	4937.000000	4937.000000	4937.000000	4937.000000	4937.000000

mean	2499.551347	49.846465	849.752279	302.580514	1.994734
std	1444.069407	29.061131	86.840313	160.804912	0.982821
min	1.000000	1.000000	700.000000	90.000000	1.000000
25%	1248.000000	24.000000	775.000000	163.000000	1.000000
50%	2497.000000	50.000000	850.000000	284.000000	2.000000
75%	3751.000000	74.000000	925.000000	387.000000	3.000000
max	5000.000000	100.000000	999.000000	1760.000000	8.000000

Here the AOV=\$302.58 which looks like a realistic value for the avg order amt placed by each customer. Hence this would be a better solution.

**b. What metric would you report for this dataset?** We can check two main metrics here:

- High Valued Customers (HVA): User\_id vs order amount and total items(To see who purchased - ?)
- High Valued shops (HVS): Shop\_id vs order amount and total items to know which shop made the

```
[37]: df_grouped=df.groupby(by='user_id').agg({'order_amount':'sum', 'total_items':
      ↪ 'sum'}).reset_index()
```

```
[38]: df_grouped
```

```
[38]:
```

	user_id	order_amount	total_items
0	607	11968000	34000
1	700	4790	30
2	701	5162	34
3	702	5286	34
4	703	6091	41
..	...	...	...
296	995	5316	35
297	996	5312	37
298	997	29236	24
299	998	2620	18
300	999	7195	47

[301 rows x 3 columns]

```
[39]: # Now lets sort
df_grouped.sort_values(by=['order_amount', 'total_items'], ascending=False)
```

```
[39]:
```

	user_id	order_amount	total_items
0	607	11968000	34000
179	878	156936	24
135	834	108342	38
88	787	85707	57
270	969	84269	50
..	...	...	...
51	750	2359	15

18	717	2337	16
20	719	2314	15
240	939	2196	15
203	902	2108	14

[301 rows x 3 columns]

- User IDs **607** is makinng bulk purchases of items and gives the max revenue (they could be a whole sale retailer or small business).
- Followed by are user ids **878, 834** who are elite buyers purchasing expensive sneakers.

These customers could be targetted for promotional offers, to encourage their purchasing capabilities and build more trust with the company

```
[29]: df_groupshop=df.groupby(by='shop_id').agg({'order_amount':'sum', 'total_items':
↪ 'sum'}).reset_index()
df_groupshop.sort_values(by=['order_amount', 'total_items'], ascending=False)
```

```
[29]:
```

	shop_id	order_amount	total_items
41	42	11990176	34063
77	78	2263800	88
88	89	23128	118
80	81	22656	128
5	6	22627	121
..	...	...	...
1	2	9588	102
99	100	8547	77
55	56	8073	69
31	32	7979	79
91	92	6840	76

[100 rows x 3 columns]

- Shop IDs **42** is doing good with a revenue of 11.9M and total item sale of 34K shoes
- Followed by are user ids **78** which sells expensive sneakers for a high revenue of 2M

These shops can be improved with more investments and seasonal offers

**0.2 Question 2:** For this question you'll need to use SQL. Follow this link to access the data set required for the challenge. Please use queries to answer the following questions. Paste your queries along with your final numerical answers below.

**0.2.1 a. How many orders were shipped by Speedy Express in total?**

```
SELECT count(OrderID) as Total_Orders
FROM Orders as O
inner join Shippers as S
on S.ShipperID=O.ShipperID
where S.ShipperName='Speedy Express'
```

ANS: 54

**0.2.2 b. What is the last name of the employee with the most orders?**

```
SELECT E.LastName, count(OrderID) as Most_Orders
FROM Employees as E
inner join Orders as O
on E.EmployeeID=O.EmployeeID
group by E.LastName
order by count(OrderID) desc
```

ANS: Peacock 40

**0.2.3 c. What product was ordered the most by customers in Germany?**

```
SELECT P.ProductName, sum(OD.Quantity) as Most_Ordered
FROM Customers as C, Orders as O, OrderDetails as OD, Products as P
where C.Country='Germany'
and C.CustomerID=O.CustomerID
and O.OrderID=OD.OrderID
and OD.ProductID=P.ProductID
group by P.ProductName
order by sum(OD.Quantity) desc
```

ANS: Boston Crab Meat

**0.2.4 Thank you**